




Visualizing Page Tables

... for Local Exploitation: Hacking Like in the Movies

Alexandru Radocea

- Developer at CrowdStrike, Inc.
 - iOS internals fan
 - Recovering software security assessor
 - Likes bringing pain to the adversary
- @defendtheworld on Twitter



Alexandru Ionescu

- Reverse Engineered Windows kernel since 1999
 - Lead kernel developer at ReactOS
- Coauthor of Windows Internals
- Founded Winsider Seminars & Solutions Inc. to provide services and training to various organizations
- Interned at Apple for a bunch of years
- Now chief architect at CrowdStrike

Georg Wicherski



- Researcher at CrowdStrike, Inc.
 - x86 & ARM low-level stuff
 - Reverse Engineering, Malware analysis
 - Exploitation and Mitigation research
- @ochsff on Twitter
- <http://blog.oxff.net/>



Introduction

Paging 101

- Translation from virtual addresses to physical
 - Virtual address: the pointers the CPU works with
 - Physical address: the actual address of a memory cell in the physical RAM chip, only used by MMU and DMA
- Virtual address unique per virtual memory space
 - Usually means per process for usermode, one shared kernel space for all processes

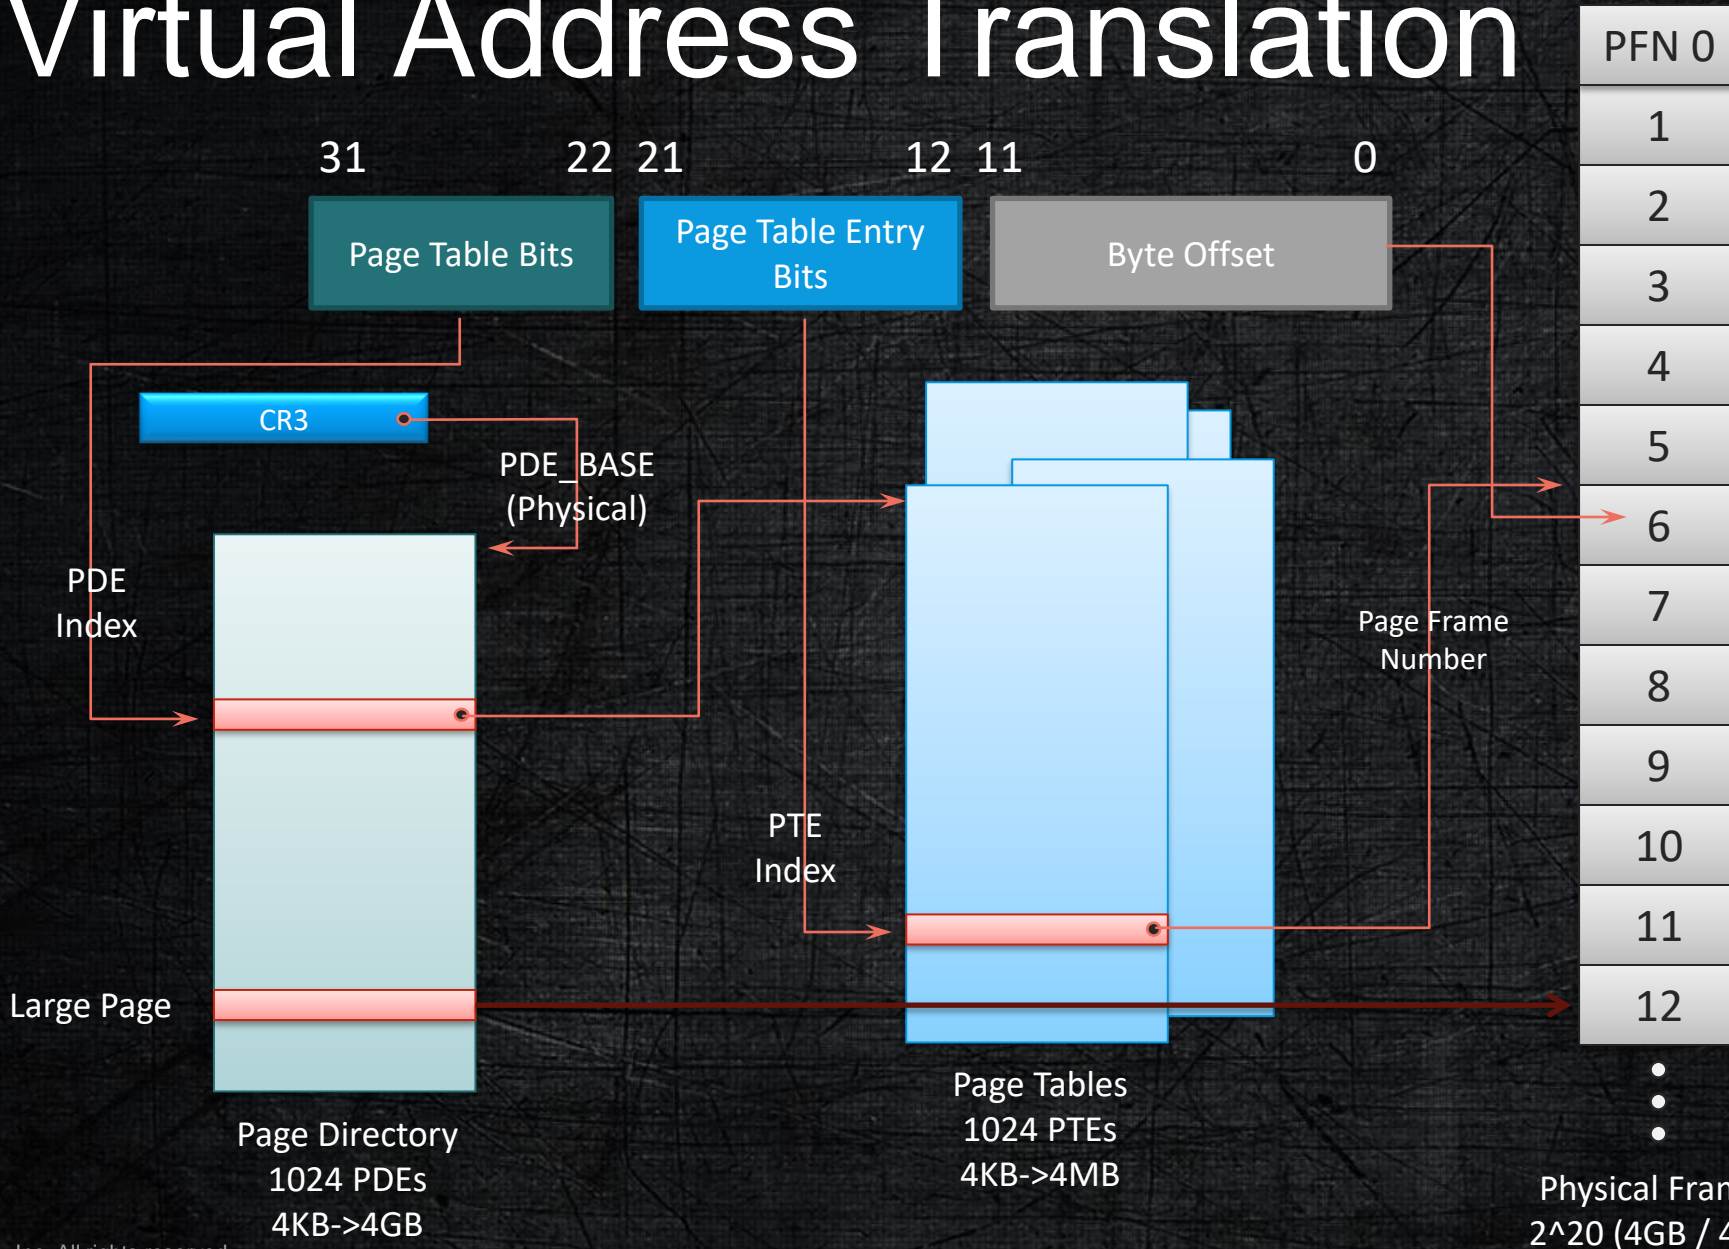
Efficient Hardware Implementation

- Group addresses into pages: block of addresses that are translated in the same way
- Cache translation results: TLB
- Hierarchical translation tables (trees) to conserve memory
 - Two levels on x86, three levels on PAE, four on x86_64
 - Two levels on ARMv7-A, three levels with LPAE

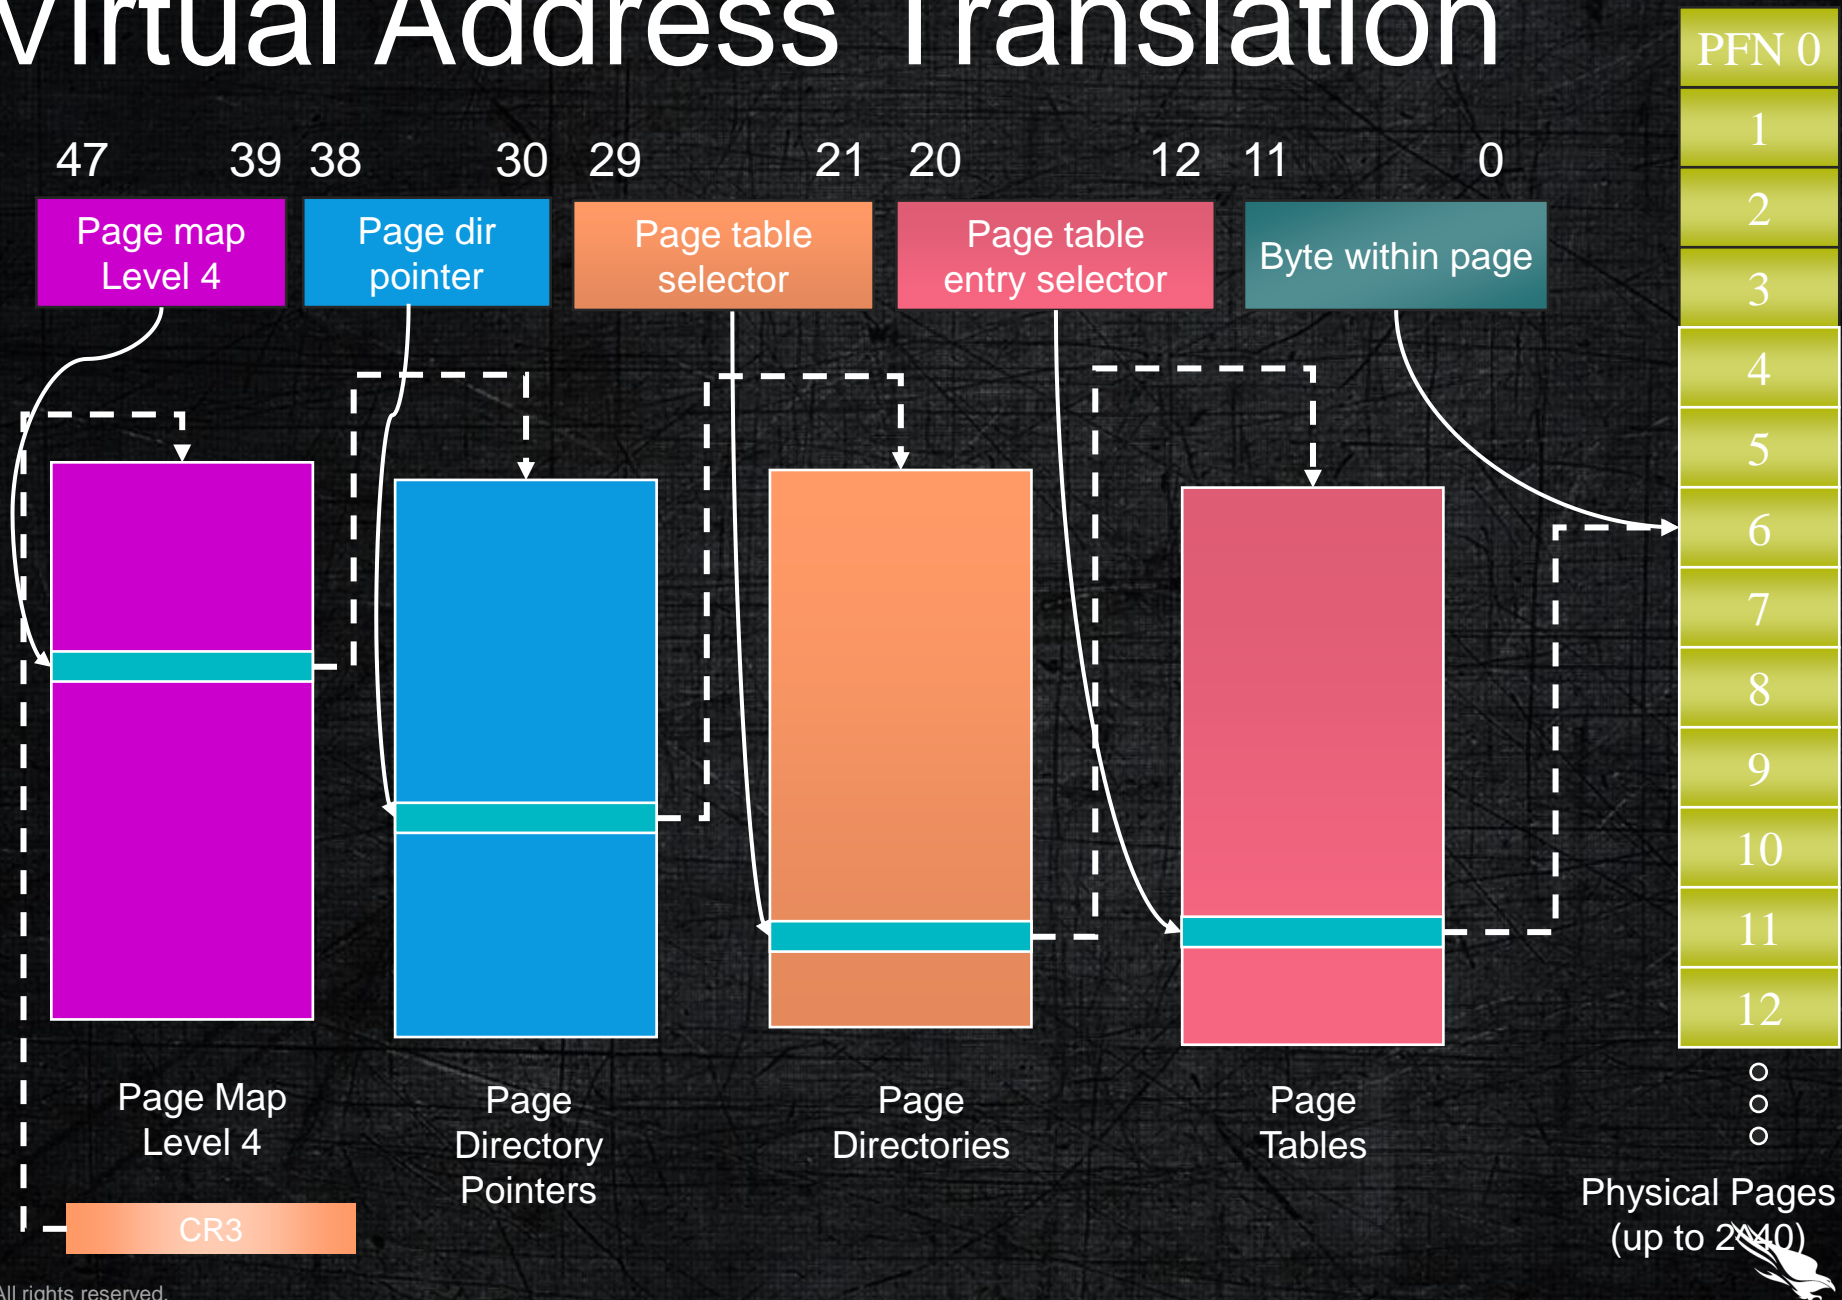
Memory Protections

- Memory protections implemented on top of paging
 - Read-only vs. read-write memory areas
 - Executable vs. data-only memory areas
 - x86_64: NX (No-eXecute) bit per page; SMEP
 - ARM: XN (eXecute-Never) bit per page; PXN
 - Privilege level to access page
 - ARM: Supervisor bit, Domains, different table sets
 - x86: Supervisor bit (CPL, SMAP)

x86 Virtual Address Translation



x64 Virtual Address Translation



What a Movie Hacker Looks for

- Mappings at repeatedly constant addresses
 - Constant physical address: Subject to reliable external physical attacks
 - Constant virtual address: ASLR bypass
- Mappings with unexpected protections
 - Read-write but not NX/XN: Classical copy shellcode and execute scenario
 - Driver specific weirdness (DMA memory, ...)



Background and Methodology


Data Collection

- Android: Both custom kernel and local exploit
- iOS: Custom driver for jailbroken device
- x86_64 Linux: Custom kernel module
- x86_64 OS X: Custom kernel extension
- Windows RT (ARM): Crash dumps
- Windows 8 x86_64: Crash dumps & custom driver
- Windows 8.1, x86_64: Crash dumps

Hilbert Curve

- Space Filling Curve
 - Well known from visualizing IP address space
 - Adapted to visualize memory address space
- Adapted scurve implementation
 - Aldo Cortesi <http://corte.si/>
 - Simple glue code to parse data collection output

Hilbert Curve Legend



User RO

Kernel RO

User RW

Kernel RW



User RX

Kernel RX

User RWX

Kernel RWX



Case Studies

Android




Android Process Comparison


1. init
2. dhcpcd
3. zygote
4. com.android.email
5. sandboxed_process0 (Chrome)




Galaxy Nexus, Android 4.2.2



Nexus 7, Android 4.2.2



Galaxy S4, Android 4.2.2 (MSM)



Android Observations

- Fixed r-x mapping at 0xffff0000 in all processes
 - 0xffff0000 is the ARM exception vectors base address
 - Abused in a *vsyscall* like manner by Linux on ARM
- Kernel .text is rwx on almost all kernels
 - CONFIG_DEBUG_RODATA not set in kernel configs
 - 3.4.x MSM kernel has RO .text
 - CONFIG_STRICT_MEMORY_RWX (Qualcomm)
 - Still has two rwx supervisor sections (1Mb pages)

Android 4.2.2 4.3 ASLR Bypass

- `__kuser_cmpxchg`: @ 0xffff0fc0
 - `arch/arm/kernel/entry-armv.S`
 - `iff *r2 == r0: *r2 := r1`
 - Brute force addresses by invoking a loop, `r0-r2` are legitimate register parameters
 - Jump past equality check for arbitrary write gadget
- `__kuser_cmpxchg64`: @ 0xffff0f60
- `ffff0008: ldr pc, [pc, #1072] ; 0xffff0440`
 - This leaks the kernel's system call handler address to user-space

Responsible Disclosure for ARM/Linux

- Initially disclosed to Google (Android)
- Quick vectors patch from Russel King, ARM
 - Randomize some code locations in page
 - Only partial solution for user-space helpers
 - Actual vector handlers are branch into adjacent page
 - Fill page with undefined instruction (ARM and Thumb)
- Randomization lacks available entropy at boot
 - Issue being discussed with security@kernel.org

Surface RT



Evolution of Windows Kernel RWX

	x86 (PAE)		x64		ARM
	Win7	Win8	Win7	Win8	Win8
Paged pool	X	X	NX	NX	NX
Non-paged pool	X	X	X	X	X
Non-paged pool (NX)	N/A	NX	N/A	NX	NX
Session pool	X	X	NX	NX	NX
Image data sections	X	X	NX	NX	NX
Kernel stacks	NX	NX	NX	NX	NX
Idle/DPC/Initial stacks	X	NX	X	NX	NX
Page table pages	X	NX	X	NX	NX
PFN database	X	NX	X	NX	NX
System cache	X	NX	X	NX	NX
Shared user data	X	NX	X	NX	NX
HAL heap	X	NX	X	NX	NX

Surface RT


- Runs Windows RT
 - Windows 8 32-bit ARM Kernel
- Locked-down kernel and user-mode signing
 - See “*Windows 8 Mitigations & ARM*” talk at Breakpoint 2012 for more low-level details...
- Suffers from 32-bit address space limitations, but has more x64 mitigations turned on
- Quad-Core Cortex A9, 2GB RAM, 32-128GB Flash

Windows RT & ARM VMSA

- Windows RT uses a VMSA configuration which brings it closest to the x86/x64 model
 - TEX is used to provide 2 software bits
 - AFE is used to provide simplified access model (hardware access bit)
- LPAE and PXN do not appear to be used
- Windows RT PTE format:
 - [0] - XN, [1] - Valid, [2,3] - Caching, [4] - Accessed
 - [5] - Owner, [6] - TEX, [7] - Writable, [8] - Copy-on-Write
 - [9] - !Dirty, [10] - Large, [11] = NonGlobal, [12-31] - Page


Surface RT Layout

- TTBR0 only
 - TTBR1 not in use
- Standard 32-bit Windows split
 - 0x00000000->0x7FFFFFFF U
 - 0x80000000->0xFFFFFFFF K
- Physical memory starts at 0x80000000
- 256 ASIDs used for TLB perf



Kernel Observations

- Kernel ASLR has difficulties due to limited address space
- Limited Kernel W^X
- Vector Page is RWX (kind of)
- I/O mappings are RWX
- Very similar layout to x86, but less RWX pages due to NonPagedPoolNx being default



Two devices side by side: Virtual view




Two devices side by side: Virtual view

- Fixed IO mappings (*MmMapIoSpace*) and usually RWX
 - Locality: *MmMapIoSpace* does not have randomization enabled, and I/O mappings are usually done quite early at boot
 - Permissions: Windows 8 introduced hacky (backward-compatible) way of asking for W^X I/O Mappings, likely not used across the board
- Fixed kernel code and RWX kernel memory was seen
 - We don't believe fixed data was seen, however (TBD)
- Userland ASLR seems very strong: nothing static seen


Where to “snap-in” a kernel payload

- Local exploits
 - No dereference protection
- Remote exploits
 - Fixed, RWX memory
 - Vector page and plenty more




ARM Vector page

- Handles interrupts
 - 0x00000000 or 0xFFFF0000
 - No requirement for RWX
- Windows has a mapping there
 - Kernel Read Execute (no Write)
- The region is all zeroes however...
- Windows uses VBAR to randomize vector table 😊





Two devices side by side: Physical view



Internet Explorer

- No obvious RWX JIT region





Windows® 8

64-bit Kernel Address Space

Observations



64-bit Kernel Address Space

Observations



The world's most advanced mobile operating system.

Dumping iOS

- Pictures depict samples from an iPhone 4S
- Injected kernel code
- Modified evas10n exploit to undo pagetable tomfoolery

iOS Layout

- TTBR Split depends on device RAM
- TTBR0 swapping on context switch
 - No shared address space

iOS 6 Security Properties

- Userland
 - Per-boot randomization (shared cache)
 - Per-execution randomization (dyld, .text, stack, heap)
 - Heap and stack separately randomized
 - W^X

iOS 6 Vector page

- Execute supervisor, read only
- No obviously useful ROP primitives




```
00000000 18 f0 8f e2 24 f0 8f e2 30 f0 8f e2 3c f0 8f e2 |....$....0...<...|
00000010 48 f0 8f e2 54 f0 8f e2 60 f0 8f e2 09 f0 a0 e1 |H...T...`.....|
00000020 fe ff ff ea 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
00000030 90 df 1d ee c0 d4 9d e5 ac d0 9d e5 04 f0 9d e5 |.....|
00000040 90 df 1d ee c0 d4 9d e5 ac d0 9d e5 08 f0 9d e5 |.....|
00000050 90 df 1d ee c0 d4 9d e5 ac d0 9d e5 0c f0 9d e5 |.....|
00000060 90 df 1d ee c0 d4 9d e5 ac d0 9d e5 10 f0 9d e5 |.....|
00000070 90 df 1d ee c0 d4 9d e5 ac d0 9d e5 14 f0 9d e5 |.....|
00000080 90 df 1d ee c0 d4 9d e5 ac d0 9d e5 18 f0 9d e5 |.....|
00000090 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
*
00000100 4f 63 74 6f 70 75 73 20 00 00 00 00 00 00 00 00 |Octopus .....|
00000110 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
*
00001000
```



iOS 6 0xffff1000

- User readable, supervisor writable page adjacent to vector table
- Available in every process
- Time, performance counters (sidechannel?)

KASLR: Kernel Code Peekaboo



Physical View (Three Captures)




iOS: Example process (MobileSafari)

- JIT (1.2MB)




iOS: Example process (MobileSafari)

- Main executable




iOS: Example process (MobileSafari)

- Shared cache




iOS: Example process (MobileSafari)

- Heap



MobileMail vs MobileSafari



iOS 6 Observations

- Evasi0n jailbreak leaves kernel mappings as RWX
- Fixed physical memory mappings across boots
 - Theoretically aids invasive attacks
 - May be useful to exploit DMA flaws from device hardware




OS X Mountain Lion

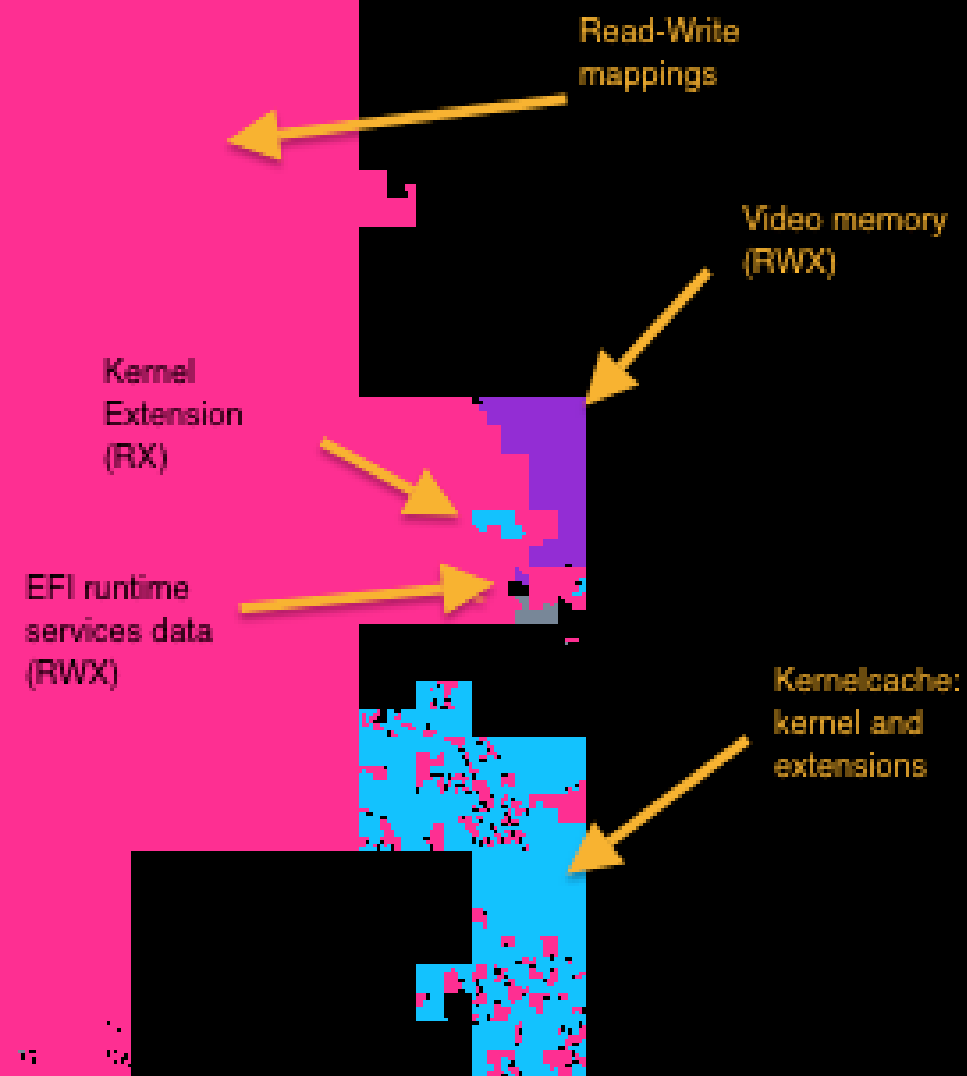

OS X Observations

- Userland
 - Per-boot randomization (shared cache)
 - Per-execution randomization (dyld, pfz, commpage, stack, heap)
 - W^X not enforced

OS X Observations

- Kernel
 - KASLR
 - Incomplete W^X
 - Unpredictable RWX regions
 - Shared address space
 - SMEP available
 - Physical addresses also randomized





Special Thanks

- Shawn Denbow for Surface RT dumps
- Evad3rs & iPhone Dev Team for iOS jailbreaks
 - <http://evasi0n.com/>
 - <http://blog.iphone-dev.org>
- Aldo Cortesi for binvis
 - <http://corte.si/posts/visualisation/binvis/>



CROWDSTRIKE