

INSECTION

AWESOMELY EXPLOITING SHARED MEMORY OBJECTS

INFILTRATE 2015

ALEX IONESCU

@AIONESCU

WHO AM I?

Chief Architect at CrowdStrike, a security startup

Previously worked at Apple on iOS Core Platform Team

Co-author of *Windows Internals 5th and 6th Editions*

Reverse engineering NT since 2000 – main kernel developer of ReactOS

Instructor of worldwide Windows Internals classes

Conference speaking:

- Infiltrate 2015
- Blackhat 2015, 2013, 2008
- SyScan 2015-2012, NoSuchCon 2014-2013, Breakpoint 2012
- Recon 2014-2010, 2006

For more info, see www.alex-ionescu.com

WHAT THIS TALK IS ABOUT

Some Windows Internals, of course ;-)

- Sections, Control Areas, Segments, Address Windowing Extensions (AWE)

Section Object Internals

- What they are
- Different types of shared objects
- How to abuse named section objects

Determining Section Ownership

- Object Type Creator Tracing
- Segment Tracking

Epic WinDBG Scripting!

Fuzzing for Section Vulnerabilities

- Three real-life case studies on my own laptop (LIVE!)

WHAT THIS TALK IS ALSO ABOUT

KASLR Bypassing Using Windows' Self-Referencing PTE Flat Mapping

- Builds upon previous research, with some new twists
- [ref: <https://labs.mwrinfosecurity.com/blog/2014/08/15/windows-8-kernel-memory-protections-bypass/> by Jeremy Fetiveau (@__x86) from MWR]
- [ref: <http://hypervsir.blogspot.com/2014/11/page-structure-table-corruption-attacks.html> by Kevin Bing (@imsky) from Intel]
- [ref: <https://www.syscan.org/index.php/download/get/5ba020d2e4b4817cb1086a0fcc2cea54/SyScan15%20Peter%20Hlavaty%20-%20Back%20To%20The%20CORE.pdf> by Peter Hlavaty (@zer0mem) from KeenTeam]

SMEP Bypassing Using AWE

- Two unpublished techniques!

Zero Day Kernel Code Signing Bypass

- This is an OFFENSIVE con, right?

PART 1 (User) OUTLINE

Windows Internals Primer

- Windows Shared Memory Model
 - Private vs Shared
 - Four Types of Shared Memory
- Windows Object Management
 - Naming and Security
- Going Deep in Sections
 - Control Areas, Segments and Subsections
 - Key Debugger Data

Hunting For Weakly Secured Section Objects

- Determining Ownership
- Fuzzing Techniques

Demo Time [Three Case Studies]

PART 2 (Kernel) OUTLINE

Tracking Kernel vs User Section Ownership

- Process Attachment Artifacts
- Kernel Object Header Flag

Difficulties of Exploiting Kernel Sections and Workarounds

- SMEP/KASLR
- Write-What-Where Conversion
- Self-Referencing PTE Space
- Address Windowing Extensions

Demo Time

QA & Wrap-Up

Windows Memory Basics

Private vs Shared Memory

Windows separates memory into **allocated** or **free** virtual address space

Allocated memory is separated into **committed** and **reserved**

Committed memory is further sub-divided into **private** and **shared**

- **Private:** new/malloc/HeapAlloc/VirtualAlloc/Stack/TEB/PEB/...
- **Shared:** Typically memory-mapped files (DLLs and EXEs)

Allocated memory is represented by a **Virtual Address Descriptor (VAD)**

Shared memory is also represented by a **Section Object**

There are multiple types and usages for shared memory...

Four Types of Shared Memory

Memory Mapped File Sections (UNNAMED, FILE-BACKED)

- Typically associated with dynamic libraries and executables from disk
 - LoadLibrary(), etc...
 - Windows calls these **Image Sections**
- But can also be data files mapped from disk
 - Windows calls these **Data Sections**
- Size comes from the file
- Page protections come from caller (**Data**) or from PE Sections (**Image**)
- They don't need a name – because they are bound by a file object on disk
 - Cache Manager and Memory Manager take care of finding the section object for an already-mapped file
- **But they CAN actually have a name – this is how KnownDLLs work**
 - Abused by myself at SyScan 2013

Four Types of Shared Memory

ALPC Sections

(UNNAMED, PAGEFILE-BACKED)

- Used and managed by Advanced Local-Procedure Call kernel IPC mechanism
 - Leveraged by RPC and DCOM, as well as many other internal services
- Size determined by ALPC attribute during ALPC Message Passing
- Protection is either **RO** or **RW**, managed by ALPC Subsystem
 - Protection changes are protected through **SEC_NO_CHANGE** and **MmSecureVirtualMemory**, **MmSecureVirtualMemoryAgainstWrites**
- They don't need a name because ALPC maps the server/client ends as needed during message passing
 - No way to get access to underlying object
- Abused by myself at SyScan 2014

Four Types of Shared Memory

Non-ALPC Sections (UNNAMED, PAGEFILE-BACKED)

- Pretty much the same as the ALPC Sections, but not managed through ALPC
 - Object Manager returns **handle** to the Section Object to the caller
- Caller must find a way to pass the handle to the process it wants to share the data with
 - **Inheritance:** Caller spawns a new child process, and the handles are inherited
 - **Forking:** Caller uses **fork()** API (HI BEN!!!), and the memory address space is inherited
 - **Duplication:** Caller uses **DuplicateHandle()** API, and the handle is duplicated to another process
- Unnamed Section Objects **do not have security checks enabled** (no ACLs)
 - Not an issue for ALPC Sections, because user-mode code doesn't have the **handle**
 - But big problem for non-ALPC Sections
 - [ref: <http://googleprojectzero.blogspot.com/2014/10/did-man-with-no-name-feel-insecure.html> James Forshaw (@tiranid) from Google]

Four Types of Shared Memory

Named Sections (NAMED, PAGEFILE-BACKED)

- Very similar to the previous type (identical at the memory manager level)
- Caller uses the Object Manager to assign a **name** to the section object
- Size and protection both specified by the creator
 - Interesting side-effect w.r.t memory consumption – creating but not mapping leaks memory!
- Any process on the system can now attempt to **open** the section by **name**
 - ACLs are critical to ensure only designated processes receive a handle to the section
 - ACLs can enforce **READ** vs **WRITE** permissions on the shared memory itself
 - Windows ACL Rule: No DACL (NULL) means **everyone has access**
- Nobody has yet systematically abused these yet!

Which to Use?

If sharing with only one process, consider **COM/RPC** (ALPC) if the communication model makes sense

- Takes care of the **handle** passing, **RO** vs **RW** semantics, and **synchronization**

Otherwise, consider security boundary needs and use **Non-ALPC Sections**

- If child/other process is expected to have **limited** rights, **do not use!**

If separate rights are needed, consider a **named section**

- Make sure to **secure** it!
- Incorrectly secured object can lead to **arbitrary** processes having **RW** access to **privileged** processes

“Objects”

Windows kernel tracks all of its shared, user-accessible resources, through a generalized Object Manager mechanism

- Also used for some kernel-internal resources (easy abstraction, so why not)
- Other kernel-mode components implement similar ideas too (Window Mgr.)

Handles **allocation** and **destruction** (with a simple reference-based GC)

Implements naming scheme (allows **hierarchical namespace** and lookup)

Leverages Security Reference Monitor (SRM) for both **discretionary access control** and **mandatory integrity control**

- Uses familiar ACL model just like for Files / Registry Keys

Exposes underlying kernel object through a per-process **handle** (index)

- Handle will **cache** the requested access rights at **creation time** and reuse them if granted

Section Objects

The **Section Object** represents the runtime semantics of a piece of shared memory

- Virtual Address, Flags (**Image** vs **Data**), Size, Session ID, Initial Protection

The actual memory backing semantics are actually held by a **Control Area**

- **Multiple section** objects can point to the **same control area** (multiple mappings of the same file)
- **Multiple section** objects can exist for a **file** (mapped once as data, and again as image)
- **Multiple control areas** can exist for the **same section** object (ask me offline if you really care how this can happen)

The **Control Area** points to a **Segment** and is made up of **Subsections**

Segment and Subsections

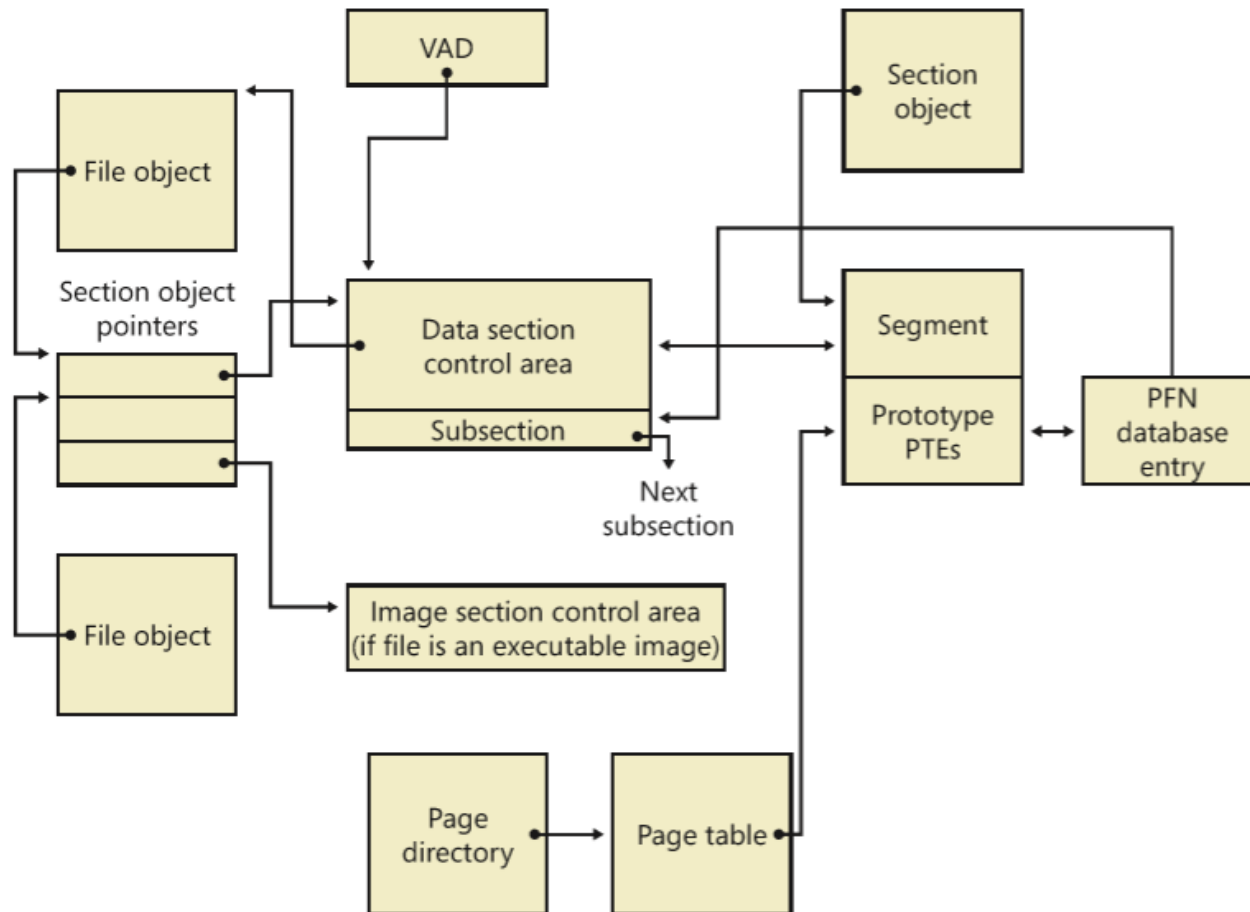
Subsections follow the **Control Area** and describe different pieces of the mapped file

- For **Image** file-backed sections, the subsections are created based on the PE sections
- For **Data** file-backed sections, the subsections are created if multiple sparse offsets of the file are mapped with different protections
- For pagefile-backed (i.e.: not file-backed) sections, there is only one **subsection**

The **Segment** describes additional internal memory-manager flags for the **Control Area**

- Points to the **Creator Process** that owns the pagefile-backed section
- Also stores address of the **first** virtual address **mapping** (if any)
- Stores the “**Prototype PTEs**” which define up the physical location of the data

So... you got all that, right?



Hunting For Insecure Sections

Finding (Un)Secured Objects

SysInternals has two great tools for looking at the Object Manager namespace

WinObj is a GUI Tool. Double-clicking will display a **properties** dialog, with a GUI ACL **Editor**

- However, it doesn't for objects that Mark hasn't written custom code for – cannot show ACLs for arbitrary objects (works for Sections, however)
- **Hard to automate**

AccesChk is a CLI Tool. Can **dump** a particular object, or **recursively** dump an object directory.

- Understands all object types
- Has flags to filter by object type, and output is **easy** to send to a **parser**

Quick Demo...

To dump all named section object ACLs, one can do

- **accesschk -s -o -t Section > acl.txt**
 - Parse with some tool to figure out what's unsecured

Can also just do quick checks:

- **accesschk -s -o -t Section | findstr Everyone**

```
\BaseNamedObjects\__ComCatalogCache__
  Type: Section
  R   Everyone
  RW NT AUTHORITY\SYSTEM
  RW BUILTIN\Administrators
  R   APPLICATION PACKAGE AUTHORITY\ALL APPLICATION PACKAGES
\BaseNamedObjects\GDA: ESENT Performance Data Schema Version 265
  Type: Section
  RW NT AUTHORITY\SYSTEM
  RW BUILTIN\Administrators
  RW NT AUTHORITY\SERVICE
  RW BUILTIN\Performance Monitor Users
  RW BUILTIN\Performance Log Users
\BaseNamedObjects\Wmi Provider Sub System Counters
  Type: Section
  RW NT AUTHORITY\SYSTEM
  RW NT AUTHORITY\NETWORK SERVICE
  RW NT AUTHORITY\LOCAL SERVICE
  R   Everyone
```

Issues with the Tools

Starting with Windows Vista, the kernel has **private namespaces**

Documented on MSDN, can be used by user-mode applications to **hide** their objects

- Won't show up in any user-mode API
- Won't show up in !object in the Debugger
 - With private symbols + an undocumented flag, you can see them though ;-)
- These aren't really any more protected than normal objects, they're just hidden

Also, the tools only show you the security descriptor, but not who **owns** the object, who **created** it, who is **accessing** it, etc..

- All this information is in the kernel debugger, however...

Object Creator Type Tracking

Windows has a poorly documented internal flag which changes the Object Manager behavior in three ways:

- It stores the **creator process** (EPROCESS Object) for each object
- It **links all the objects** of the same type together (no need for memory scanning or fingerprint)
- It (optionally) takes a complete **stack backtrace** at object creation time

You can enable this really useful behavior with the **Global Flags Utility**

- Set the option “**maintain a list of objects for each type**”
- Can also use **!gflag +otl** in the kernel debugger (sets flag 0x4000)

This is exposed in two ways

- A new **NtQueryObject** information class is exposed, dumps all objects of a given type
- A new capability is added to the **!object** WinDBG extension to do the above too

Dump all the things!

Once the Global Flag is set, you can use the following syntax in WinDBG:

- **!object 0 Section**

This will dump **all** section objects

- We only care about the ones that have names...
- No way to filter

Well... we can filter with the magic “**.foreach**” WinDBG scripting command

- Allows parsing output from the debugger, and choosing/skipping which **tokens** to use as **input**

But instead of using a WinDBG extension and then a filtering script...

- Let's just write our own script

Linked Objects

All objects of the same type have an **OBJECT_TYPE** structure which owns the **type** of the object (Process vs Thread vs Section)

When the **+otl** flag is enabled, a **TypeList** doubly-linked circular list is used to link the objects of the same type together

- `dt nt!_OBJECT_TYPE TypeList`

Every unique object has a **main header** and a possible set of **extra headers** that define the attributes for the particular object

- Name, memory quota consumption, etc...
- When **+otl** is used, a new **creator info header** is allocated for each object
- `dt nt!_OBJECT_HEADER_CREATOR_INFO`

The object creator info header is always **on top of the main header**

- `dt nt!_OBJECT_HEADER`

Enumerating a Linked List

The debugger has a **!list** extension to iterate over a linked list

Each iteration can execute a command typed between the **-x** “...” input

- The current entry being iterated is stored in a **@\$extret** register variable

How do we get the linked list for the section object?

- Each object has an exported symbolic name
- Section: **nt!MmSectionObjectType**

This will iterate over every section object:

- `r? @$t0 = *(nt!_OBJECT_TYPE**)@@(nt!MmSectionObjectType)`
- `!list -x "?? @$extret + sizeof(nt!_OBJECT_HEADER_CREATOR_INFO) + sizeof(nt!_OBJECT_HEADER)" @@(&@$t0->TypeList)`

Filtering Named Objects

If an object has a name, it has an **object name info header**

If an object has any type of **extra headers**, the **main header** describes that by setting the corresponding **bit** into the **InfoMask** bitfield

- Bit Value **0x2** is used for the name info header
- `dt nt!_OBJECT_HEADER InfoMask`

We can filter named objects by doing something like this:

- `$$ T1 is the Object Header`
`.if (@@((@$t1->InfoMask & 0x2) != 0))`
`{`
 `$$ This is a named object`
`}`

Finding the Security Descriptor

If an object is protected by a security descriptor, the **SecurityDescriptor** field in the **main header** will point to it (almost...)

- `dt nt!_OBJECT_HEADER SecurityDescriptor`

The **!sd** extension will then display the DACL and SACL in the security descriptor

- Add **flag 0x1** after the pointer and you'll get SID->Name translation too

But the pointer is actually a **fast reference**

- Bottom 3 (x86/ARM) or 4 (x64/ARM64) bits store a reference count

We must align it first:

- **\$\$ T1 is the Object Header**
`!sd @@(((unsigned int64)@$t1->SecurityDescriptor & ~0xF)) 1`

We are particularly interested in DACLs that are either **empty**, or which grant **“Everyone”** RW access

- No way to filter at the extension level, we need **.foreach** or a custom script

Finding the Owner

When we talk about the “owner” of an object, there’s really two things we could be talking about

The “Owner SID”, which is a security concept, stored in the security descriptor

- The owner **does not imply creatorship**. It is merely used to allow certain additional permissions to the SID when object access is attempted. **Anyone can set the owner to someone else when creating a new object.**

The actual “**Creating Process**” which Windows does not normally track...

- But we have the **+otl** flag!

But remember that for section objects, the **Segment** does also independently track the creator!

Creator Lookup with Header

Get the **Object Header** for the **Section**, get the **Object Creator Info**, and finally get the **creator PID**, feeding it into **!process**:

- `$$ Section Object in T0`
`!process @@(((nt!_OBJECT_HEADER_CREATOR_INFO*)((unsigned int64)#CONTAINING_RECORD(@$t0, nt!_OBJECT_HEADER, Body) - sizeof(nt!_OBJECT_HEADER_CREATOR_INFO)))) ->CreatorUniqueProcess) 0`

Or if you already have the header...

- `$$ Section Object Header in T1`
`!process @@(((nt!_OBJECT_HEADER_CREATOR_INFO*)((unsigned int64)@$t1 - sizeof(nt!_OBJECT_HEADER_CREATOR_INFO)))) ->CreatorUniqueProcess) 0`

Creator Lookup with Segment

Get the **Segment** from the **Section Object**, then dereference **u1.CreatingProcess**

- `dt nt!_SEGMENT u1.CreatingProcess`

In **Windows 10**, the **Section Object** has a pointer to the **Control Area** at **u1.ControlArea**

- `dt nt!_SECTION u1.ControlArea`

So in Windows 10:

- **\$\$ Assumes Section Object in T0**
`!process @@(@$t0->u1.ControlArea->Segment->u1.CreatingProcess) 0`

In Windows <10, the **Section Object** has a pointer to the Segment

- And the `nt!_SECTION` type is not in the symbols!
- You can use `nt!_SECTION_OBJECT`, and dereference the **Segment** field
- **CAREFUL:** Type is `nt!_SEGMENT`, not `nt!_SEGMENT_OBJECT` as shown

Security Descriptor's Owner

Get the **Owner** from the **Security Descriptor**, then add the value to the Security Descriptor pointer

- This is because it is stored in **relative** form
- `dt nt!_SECURITY_DESCRIPTOR_RELATIVE Owner`

The resulting pointer is a **SID** structure

- `dt nt!_SID`

We can use the **!sid** extension to dump it, and pass flag **0x1** to perform SID->Name conversion

- **\$\$ Section Object Header in T1**
`!sid @@(((nt!_SECURITY_DESCRIPTOR_RELATIVE*)((unsigned int64)@$t1->SecurityDescriptor & ~0xF))->Owner + ((unsigned int64)@$t1->SecurityDescriptor & ~0xF)) 1`

Might display either a **user** (Bob) or a **group** (Administrators)

Creator Process To Account

The creator process is interesting for many purposes, but doesn't quickly tell us if the section object is worth attacking

- Process could “sound” important but run with limited privileges, etc...

If we take a look at the **Token** of the Process, this will tell us the **User Account** that it's running as

- The **Token** field in **EPROCESS** stores a **fast reference** to the token object

The **!token** extension will display and format a token, and the **-n** flag will convert SIDs to names

- We're interested in SYSTEM/LOCAL SERVICE, for starters

Again, filtering is not possible through the extension, so we have to use **.foreach** or write a custom script

- The **UserAndGroups** field in the **Token** stores the SID of the user account
- **!sid @@(@\$t1->UserAndGroups[0].Sid) 1**

Good Script Starting Point...

Assuming **+otl** was enabled in the Global Flags, use **!list** to enumerate over all the section objects

For each entry:

- Filter out **named** sections only
- Filter out **Image**, **Data**, or **Physical** sections (they don't have an owner process and are not relevant)
- Print **Object Creator** PID (from **Header**)
- Print **Section Creator** PID (from **Segment**)
- Print **User Account** of Section Creator PID (from **Token**)
- Check if **DACL** is NULL (from **Security Descriptor**) and print a **WARNING**
- If **DACL** is != NULL, parse each **Access Control Entry (ACE)**
 - For each **ACE**, print the **SID**
 - Check if the **ACE** is S-1-1-0 (**Everyone**), and if the access is **SECTION_ALL_ACCESS** (0xF001F) and print a **WARNING**

Time To Make Your Eyes Bleed

```
r? @$t0 = (nt!_OBJECT_HEADER_CREATOR_INFO*)@$extret
r? @$t1 = (nt!_OBJECT_HEADER*)(@$t0 + 1)

.if (@@(@$t1->InfoMask & 0x2))
{
    r? @$t2 = (nt!_OBJECT_HEADER_NAME_INFO*)@$t0 - 1
    r? @$t4 = (nt!_SECURITY_DESCRIPTOR_RELATIVE*)((unsigned int64)@$t1->SecurityDescriptor & ~0xF)
    r? @$t5 = (nt!_SECTION*)&@$t1->Body

    .if (@@((@$t5->u1.ControlArea->u.Flags.Image == 0) && \
        (@$t5->u1.ControlArea->u.Flags.File == 0) && \
        (@$t5->u1.ControlArea->u.Flags.PhysicalMemory == 0)))
    {
        r? @$t6 = @$t5->u1.ControlArea->Segment->u1.CreatingProcess
        r? @$t7 = (nt!_TOKEN*)(@$t6->Token.Value & ~0xF)

        .printf /D "\nOBJECT: %p [<b>%msu</b>] K: %d\n", @$t1, @@(&@$t2->Name), @@(&@$t1->KernelObject)
        .printf /D "\tPID: <b>%lx</b> EPROCESS_PID: <b>%lx</b>\n", @@(&@$t0->CreatorUniqueProcess),
        .printf /D "\tOWNER "
        !sid @@((unsigned int64)@$t4 + @$t4->Owner) 1
        .printf /D "\tTOKEN\t"
        !sid @@(&@$t7->UserAndGroups[0].Sid) 1

        .if (@@(&@$t4->Dacl == 0))
        {
            .printf /D "<col fg=\"changed\">\t<b>WARNING!!! WARNING!!! WARNING!!!</b></col>\n"
        }
        .else
        {
            r? @$t8 = (nt!_ACL*)((unsigned int64)@$t4 + @$t4->Dacl)

            $$>< c:\insection\aclcheck.txt
        }
    }
}
```

Live Results...

```
OBJECT: fffffc0016fbac5a0 [RTKAPO_NSM_AEC{d9945156-9691-442f-99a2-1b94e0005708}]
PID: 1024 EPROCESS_PID: 1024
OWNER SID is: S-1-5-19 (Well Known Group: NT AUTHORITY\LOCAL SERVICE)
TOKEN SID is: S-1-5-19 (Well Known Group: NT AUTHORITY\LOCAL SERVICE)
DACL SID is: S-1-1-0 (Well Known Group: localhost\Everyone) with ACCESS: f001f
WARNING!!! WARNING!!! WARNING!!!
```

```
OBJECT: fffffc0016d9a24e0 [FloresvilleIWMSInterface_Vista]
PID: 6cc EPROCESS_PID: 6cc
OWNER SID is: S-1-5-18 (Well Known Group: NT AUTHORITY\SYSTEM)
TOKEN SID is: S-1-5-18 (Well Known Group: NT AUTHORITY\SYSTEM)
WARNING!!! WARNING!!! WARNING!!!
```

```
OBJECT: fffffc0016dadeed0 [FoxitCloudUpdateSvcShareMemoryForPhantomPDF]
PID: 878 EPROCESS_PID: 878
OWNER SID is: S-1-5-32-544 (Alias: BUILTIN\Administrators)
TOKEN SID is: S-1-5-18 (Well Known Group: NT AUTHORITY\SYSTEM)
DACL SID is: S-1-1-0 (Well Known Group: localhost\Everyone) with ACCESS: f001f
WARNING!!! WARNING!!! WARNING!!!
```

```
OBJECT: fffffc0016ef03570 [S-1-5-21-3704014639-4232294138-2895164880-1001UserDataHostPidMemory]
PID: 1150 EPROCESS_PID: 1150
OWNER SID is: S-1-5-21-3704014639-4232294138-2895164880-1001 (User: SAMMY-PC\Alex Ionescu)
TOKEN SID is: S-1-5-21-3704014639-4232294138-2895164880-1001 (User: SAMMY-PC\Alex Ionescu)
DACL SID is: S-1-5-18 (Well Known Group: NT AUTHORITY\SYSTEM) with ACCESS: f001f
DACL SID is: S-1-1-0 (Well Known Group: localhost\Everyone) with ACCESS: f001f
WARNING!!! WARNING!!! WARNING!!!
DACL SID is: S-1-15-2-1 (Well Known Group: APPLICATION PACKAGE AUTHORITY\ALL APPLICATION PACKAGES) with ACCESS: f001f
```

Fuzzing Sections

Using WinDBG for Analysis

Easiest way (if it works), is to use the local kernel debugger (we have access to all processes this way) and **fill** the section object's memory with **garbage/41414141/etc**

First, we have to **attach** to the right process:

- `.process /P <creator process>`
- `.reload /user`

Then **dump** the section memory to see what's there:

- `dc <Segment->u2.FirstMappedVa>`

Also use the **!pte** extension to see the **page permissions**

- Bi-winning if the memory is **executable**

This might not work if the creator never **mapped** the section...

- Or has since **unmapped** it (and maybe mapped **elsewhere**)

Using WinDBG for Fuzzing

We can use the fill (**f**) command in WinDBG to write some garbage...

- **f <virtual address> 41 L1000** to write a page's worth of 414141...

This **won't work** from the local kernel debugger! ☹️

- It will work if you're remotely attached, or if you're in user-mode

Workaround is to use **physical** memory address

- Read the page frame number from **!pte** and multiply by **@\$pagesize**
- Then use **fp** instead of **f**

This will work and since we're on a live system TLB flush will happen on its own anyway

Check for any **changes** to the process

- Strange memory consumption, **crash**, etc
- Having user-mode WinDBG attached can easily display **exceptions** / DbgPrint

Caveat Emptor

Do this on most processes and absolutely **nothing** will happen

Put yourself in the software engineer's shoes

- How does the **reader** know that the memory has been **written** to (Hint: the answer is not “hardware memory breakpoints”)?
- Likewise, how does the **writer** know that it can write **more** (**reader** is **finished** reading)

Some **synchronization object** must obviously be used

- Except when ALPC is doing it for you – but these are named sections
- Could be a **mutex**, **event**, **semaphore**

If the permissions of the **synchronization object** aren't equally as weak as the section object, you may not be able to achieve **instant** exploitation...

- Or perhaps **no** synchronization object is used at all

More Complex Scenarios

No **synchronization object**?

- Maybe a **timer** is employed: every n seconds, some thread reads changes in the memory structure
- Maybe some other **external source** (I/O is completed, or **window message** is received) is causing the thread to check the memory

No access to **synchronization object**?

- Perhaps you can **re-create the scenarios** that are normally required by the privileged client to signal the object on your behalf
- This is likely going to **overwrite** your fuzzed data with authentic data, however

However, in the vast majority of cases, there's a **synchronization object**, and it's usually as secure as the section object itself

Live Case Studies

Floresville Interface

Owned by **wlanext.exe** (“Windows Wireless LAN Extension Host”)

- However, not a Windows/Microsoft bug per se – this is Intel’s Extension

Appears to contain **encrypted data**

Reverse engineering shows:

- **10 equally-sized buffers**, each one associated with its own event
- One thread **waits** on all the 10 events
- Each buffer is **~1794 bytes**, and the entire shared memory is allocated x10 that size
- Thread copies from offset `SignaledEvent*1794`, and always copies 1794 bytes
- Sends a **window message**, goes back to **waiting**

Another thread wakes up to **handle the window message**

- **Decrypts** the shared memory buffer (bails out if decryption failed)
- **Parses** it based on various encoded **commands** (each with their payload)

Floresville Interface

The screenshot displays three overlapping windows from Windows Task Manager:

- Event List:** A table showing various system events. The selected event is:

Type	Name	Handle	Granted Access
Event	\BaseNamedObjects\MUROC_ROAM_END_EVENT	0x26c	Full control
Event	\BaseNamedObjects\MUROC_DISASSOC_START_EVENT	0x270	Full control
Event	\BaseNamedObjects\MUROC_CONN_START_EVENT	0x274	Full control
Event	\BaseNamedObjects\MUROC_CONN_END_EVENT	0x278	Full control
Event	\BaseNamedObjects\MUROC_NETWORK_STATE_CHANGE_E...	0x27c	Full control
Event	\BaseNamedObjects\MUROC_PEER_STATE_CHANGE_EVENT	0x280	Full control
Event	\BaseNamedObjects\MUROC_RADIO_STATE_CHANGE_EVENT	0x284	Full control
Event	\BaseNamedObjects\S24TxDataPresentEvent_Vista0	0x2d8	Full control
Event	\BaseNamedObjects\S24RxDataPresentEvent_Vista0	0x2dc	Full control
Event	\BaseNamedObjects\S24TxDataPresentEvent_Vista1	0x2e0	Full control
Event	\BaseNamedObjects\S24RxDataPresentEvent_Vista1	0x2e4	Full control
Event	\BaseNamedObjects\S24TxDataPresentEvent_Vista2	0x2e8	Full control
- Stack - thread 2740:** A list of stack frames for thread 2740. The selected frame is:

Index	Name
0	ntoskrnl.exe!KiCheckForKernelApcDelivery+0x186
1	ntoskrnl.exe!KeWaitForMultipleObjects+0x1abc
2	ntoskrnl.exe!KeWaitForMultipleObjects+0xd35
3	ntoskrnl.exe!KeWaitForMultipleObjects+0x22a
4	ntoskrnl.exe!ObWaitForMultipleObjects+0x2b5
5	ntoskrnl.exe!RtlCompareUnicodeString+0x426
6	ntoskrnl.exe!setjmpex+0x37d3
7	ntdll.dll!ZwWaitForMultipleObjects+0xa
- wlanext.exe (1664) Properties:** A window showing the properties of the process. The 'Performance' tab is selected, displaying a table of threads:

TID	CPU	Cycles D...	Start Address	Priority
3796			ccxplugin.dll+0x46b4	Normal
3788			ccxplugin.dll+0x46d0	Highest
3792			ccxplugin.dll+0x46ec	Highest
2740			iwmssvc.dll!Dot11ExtIhvInitVirtualStation+0x13820	Normal

```

1kd> ?? (nt!_OBJECT_HEADER_NAME_INFO*)((unsigned int64)#CONTAINING_RECORD
struct _OBJECT_HEADER_NAME_INFO * 0xffffe001`610e9a30
+0x000 Directory      : 0xfffffc001`0b2b5920 _OBJECT_DIRECTORY
+0x008 Name           : _UNICODE_STRING "S24TxDataPresentEvent_Vi
+0x018 ReferenceCount : 0n0
  
```

Floresville Interface

May allow for **information disclosure**

May have subtle parsing bugs

But ultimately requires a lot of **effort** to fuzz since it requires **encrypting** the payload (otherwise it gets dropped)

Let's see if we have more success with **other** sections...

PhantomPDF Interface

Owned by **FCUpdateService.exe** (“Foxit Cloud **Safe Update** Service”)

Appears to contain **JSON data** when populated, but **usually empty**

Is present even if user **unchecks** the option to install Phantom PDF (Foxit’s Cloud storage service)

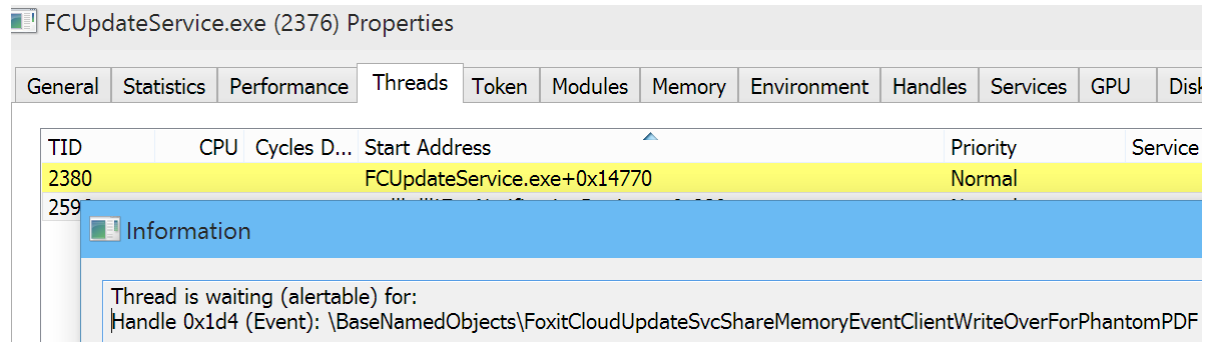
First field is the **size** of the buffer

- Directly **trusts** it and uses it during **memcpy** operations
- Also immediately **assumes** data is correctly **formatted** JSON and begins parsing it

Not only is it **trivial** to **exploit**, but I also wonder if the JSON payloads themselves can be legitimately laid out but malicious in nature

- i.e.: do a JSON payload attack, not fuzzing the JSON data/memory itself

PhantomPDF Interface



Key	HKU\DEFAULT\Software\Microsoft\Windows\CurrentVersion\Explorer	0x170
Directory	\BaseNamedObjects	0xa8
Event	\BaseNamedObjects\FoxitCloudUpdateSvcShareMemoryEventClientWriteOverForPhantomPDF	0x1d4
Event	\BaseNamedObjects\FoxitCloudUpdateSvcShareMemoryEventServerWriteOverForPhantomPDF	0x1d8
Section	\BaseNamedObjects\FoxitCloudUpdateSvcShareMemoryForPhantomPDF	0x1c8
Mutant	\BaseNamedObjects\FoxitCloudUpdateSvcShareMemoryMutexForPhantomPDF	0x1cc

```
(948.a24): Access violation - code c0000005 (first chance)
First chance exceptions are reported before any exception handling.
This exception may be expected and handled.
FCUpdateService+0x14a97:
00984a97 f3a5          rep movs dword ptr es:[edi],dword ptr [esi]
0:001:x86> db edi
018fc018  00 00 00 00 9c 00 93 00-9c 00 93 00 00 00 00 00 .....
018fc028  00 00 00 00 00 50 41 41-00 c0 41 41 d2 2f 7d c4 .....PAA..AA./}.

```

RTK APO NSM AEC Interface

Owned by **Audiodg.exe** (“Windows Audio Device Graph Endpoint”)

- However, not present by default... not a Windows component

APO = Audio Processing Objects

- **COM Interface** through which 3rd party sound card vendors can apply DSP
- Likely my **sound card manufacturer** (Hint: R***t*k)
 - Or maybe it's a Rootkit

Couldn't find an obvious **synchronization object** that it's associated with

The related threads in Audiodg.exe are each waiting on 64, 48, 36 threads, respectively (and there are two pairs of each)

So, use another technique to discover use:

- Set **hardware breakpoint** on memory and see when it's used!

RTK APO NSM AEC Interface

audiodg.exe (8588) Properties

General Statistics Performance Threads Token Modules Memory Environment Handles G

Name	Verified Signer	Description	Base Address	
RltkAPO64.dll	Realtek Semicon...	Realtek(r) LFX/GFX DSP component	0x6aa20000	3.1
ntdll.dll	Microsoft Windows	NT Layer DLL	0x7ffc12c80000	1.7
combase.dll	Microsoft Windows	Microsoft COM for Windows	0x7ffc12a20000	2.3
msvcrt.dll	Microsoft Windows	Windows NT CRT DLL	0x7ffc115c0000	62
clbcatq.dll	Microsoft Windows	COM+ Configuration Catalog	0x7ffc114b0000	65
setupapi.dll	Microsoft Windows	Windows Setup API	0x7ffc112d0000	1.7
advapi32.dll	Microsoft Windows	Advanced Windows 32 Base API	0x7ffc111e0000	66
ole32.dll	Microsoft Windows	Microsoft OLE for Windows	0x7ffc10f40000	1.2
user32.dll	Microsoft Windows	Multi-User Windows USER API Cli...	0x7ffc10dd0000	1.4

audiodg.exe (8588) (0xac6ca20000 - 0xac6ca2a...

```
00000000 01 00 00 00 fe ff 02 00 80 bb 00 00 00 dc 05 00 .....
00000010 08 00 20 00 16 00 20 00 03 00 00 00 03 00 00 00 .....
00000020 00 00 10 00 80 00 00 aa 00 38 9b 71 f7 01 00 00 .....S.q....
00000030 0a 00 00 02 00 00 00 00 00 00 00 00 10 9f 00 00 .....
00000040 b0 ce 71 cc a2 75 a0 01 e0 01 00 00 00 00 00 00 .....q.x....
00000050 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
00000060 1d 2b ad 2f 50 a7 63 2f 66 ff 9d ae 8e 5c 03 2e ...+./E.c/fo.\.
00000070 c6 d0 09 b0 21 12 8b 2f 7b af ba 2f 52 f0 a7 2e ...!./[.../R...
00000080 22 1d 16 2f e4 1b 30 af e5 91 c5 af 56 ae f5 2f .../.../...V.../
00000090 77 1a 1a 30 50 63 1a ae e4 42 be 2f 8d dd a7 2f ...0Pc...B./.../
000000a0 a6 54 d6 af 53 d1 f7 ae 99 77 03 2f e9 dd cf 2f ...T...S...W./.../
000000b0 77 5b 95 2f 2b ef c1 2f ca a8 0a 2f 11 40 f1 ae w[./+.../.../0...
000000c0 fc 31 23 b0 0e a6 af 2e 68 4c 4d 2f 51 7e 41 2e ...1#...hLM/Q~A...
000000d0 f5 1f 7c 2e 0f 91 48 2e 32 38 4e 2f cf b7 04 b0 ...H.H.28N/....
000000e0 1a 1a fd 2e 40 8a 78 2d c1 8d 02 b0 84 6a 5c 2f ...1#...x...j\...
000000f0 50 17 73 2f 84 a4 cf 2d 7a 58 3b 2f cd 9a be af ...P...2Z/...
00000100 b0 1b 33 2d bd c0 e9 ae 56 05 fe 2e 60 71 1d 2f ...3...V...q/...
00000110 01 08 e6 2f 1d c5 a7 af 6a a6 04 b0 a4 d0 2d 2e .../...m.../...
00000120 91 16 fa af 17 e2 d7 ae 5c c1 c8 2f ad d0 b8 ae .../.../.../...
00000130 94 ea a3 2f 3e d6 4b 2f c9 9c 08 b0 80 5d 4e 2f .../>.K/...]N/...
00000140 fb e0 dc 2f db a1 49 2f 75 37 48 af 84 7f 9a 2f .../...I/uTH.../
00000150 ff 14 7a 2e b4 b6 af 2e d6 41 af 82 ce 6c 2f ...2...R...1/...
00000160 00 8a d6 ae f4 05 ba af 09 c0 b1 ad 05 55 10 b0 ...U...
00000170 6b 6e 81 ab 60 75 d9 2c ec 19 cf ae 60 9f 80 2e ...x...u.../...
00000180 08 a2 a2 af e3 aa 17 b0 32 0b 83 af a4 30 ec 2d .../...2...0...-
00000190 00 b0 28 2a 02 cb 04 af fe 54 02 30 0f 14 93 2e ...(*...T.0...
000001a0 68 d7 91 2f 0e 46 e6 ae 9a 01 82 ae fa e4 e0 af ...h.../...F...
000001b0 80 98 bb 2e fb 4c f5 af 0a 2f c1 2e e4 36 da af ...L.../...6...
000001c0 52 25 45 2f 90 fd ee 2d ca f9 c0 2f 4e 93 31 ae ...R&E/.../...N.1...
000001d0 f0 f7 8d ae 3c 26 b1 af 1c d0 0e 2f e6 92 d9 2e ...<.../...
000001e0 62 be 1a af 5c 15 ab af 04 ca 39 2f be cb 5a 2f ...b.../...9.../2...
000001f0 84 c3 98 af 66 d1 a3 af 0b 7f e8 2f e5 87 e9 2f ...f.../.../
00000200 d8 eb 64 2e 36 3e 0b af 00 4d de af 68 7a 4c af ...d.6>...M.hzL...
00000210 80 37 21 2f 93 e7 13 2f b8 85 6c ad 32 31 51 ae ...7/.../...1.21Q...
00000220 cc 29 ff 2d 56 05 67 af 7f e6 86 af 17 87 0c 30 ...-V.g.../...W.0...
00000230 fd c4 d4 2f 11 43 cf 2e 4e 50 56 2f 78 de 13 b0 .../...C.NFV/x...
00000240 33 a9 0b af e7 9b af ff c7 5e ae 38 c1 30 2f ...e.../...8.0/...
00000250 e6 0b 3b 2f a7 91 99 2f 6c 13 14 af b1 b0 92 af .../.../...1...
00000260 d1 fb c7 2f 35 2c cb af 79 e0 16 ae be b4 d6 ae .../5...Y...
00000270 7f 3d 0f ae 04 f6 04 30 b0 71 20 2f ec 51 c6 ad ...=.../...0.g /Q...
00000280 d3 c4 a6 2f 9a 8a 01 2f 87 95 bf ae fa 7f 02 30 .../.../...g.../
00000290 5d 18 22 af b2 cf 14 af 87 92 94 2e 27 d0 9a 2e ...j.../.../...
000002a0 85 6d a2 af c0 3c 2f d3 21 9d 2f 8e ff 8b 2e ...m...<./.../...
000002b0 33 a2 07 ae b4 81 50 af 08 be 54 af 57 72 8e af ...3...P...T.Wr...
000002c0 c8 a3 99 ae c1 83 84 af 1a d9 34 af 9e a4 6d 2f ...4.../...
000002d0 96 1a f6 2e 17 be 94 ae cf 87 d2 af a2 ae d0 af .../...
000002e0 40 3d a4 ac 4f 30 87 af 47 d6 a8 2f 8a 1c a5 af ...=...00...G.../...
000002f0 2d 49 08 b0 97 8e 0c ad 3a eb 51 2f 05 bf b3 2f ...-I.../...Q.../...
```

Re-read Write Goto... Save... Close

Breakpoint 0 hit

msvcrt!memmove+0x1cf:

00007ffc`11633c0f 75ef jne msvcrt!memmove+

0:006> dd 00000072`26840000

00000072`26840000 ffffffff ffffffff ffffffff ffffffff

00000072`26840010 ffffffff ffffffff ffffffff ffffffff

00000072`26840020 ffffffff ffffffff ffffffff ffffffff

00000072`26840030 ffffffff ffffffff ffffffff ffffffff

RTK APO NSM AEC Interface

Ultimately was able to achieve a **crash**

- But **not** under **repeatable** circumstances

Only discovered this interface by chance on the plane on the way here...

- Doesn't show up until you actually **play music** using Windows Media Player
- Or **watch a movie**

Very interested to pursue research on this particular object, since it could allow for attack against **protected processes**

- See my blog if unfamiliar with the concept of **Protected Process Light**
- No symbols for latest Windows 10 build are making things hard
- And can't exactly run this in a virtual machine

But ultimately, we are going to look at an even **more interesting** class of section objects, so let's leave **DRM** world behind for now...

Closing Remarks on Methodology

The script currently only looks for “**Everyone**” being granted **SECTION_ALL_ACCESS** or for a **NULL DACL**

These are the most egregious ways of having insecure section objects, but not the only ones

- Granting **WRITE_DAC**, **SECTION_MAP_WRITE** is just as bad
- Granting access to the “**Users**” group is also bad
 - Or “**Remote Users**” or “**Authenticated Users**” ...
- Even granting access to the **current user** can be bad

We didn’t look for **ALL_APPLICATION_PACKAGES** either

- This is the “**Everyone**” of the **AppContainer** world (Windows 8 **sandboxes**)
- AppContainers are vulnerable to **NULL DACL**, but not to “**Everyone**” alone

With some improvements, you can probably find more vulnerabilities

Kernel Exploitation

Kernel Owned Section Objects

The API to create named section objects can also be used by **drivers**

- Will typically be owned by **SYSTEM** in the **Security Descriptor**

How can we identify such **kernel-created** section objects?

- Looking for **PID 4** (“System Process”) is one way, but the Windows Kernel runs in **arbitrary** process address spaces when syscalls or interrupts happen
- So a driver could’ve created the section object while in **another context**
- Instead, we can use a few **artifacts** of kernel-created sections...

Note in the script output, some **EPROCESS_PIDs** did not match **PIDs**

- Why would the object header have a different PID than the segment?
 - **Segment** tracks **EPROCESS**, object **header** tracks **PID**...
- Kernel code can “**attach**” to a process to access its memory, handle table, etc

PsGetCurrentProcess returns “Current” (**Attached**) Process,
PsGetCurrentProcessId returns thread’s “Owner” (**Original**) Process ID.

- **ReactOS gets this wrong!!!**

Kernel Created Objects

But the kernel doesn't necessarily have to be attached, so the artifact may not be **visible**

And the artifact can also false-positive if the EPROCESS is **re-used** after the creator process has died but the section has been **kept around** by another mapper

- **PID** of **dead** process and **PID** of **new** EPROCESS will now **mismatch**
- In pathological cases, both PID and EPROCESS will still match but point to an unrelated process

A better way is to use the Object Header's **KernelObject** flag

- Set **automatically** whenever a component in kernel-mode is creating an object

We can extend the hunting script to also **identify** kernel-created Objects

- Reveals the existence of a **\Win32kCrossSessionGlobals** structure

But only allows **SYSTEM/Administrators** to modify it. Is this a problem?

The Debate

A company spends millions of dollars, resources, and lines of code to create a system where

- Kernel code must be **signed** (by the **company itself**, in Windows 10!)
 - With a **SHA-256 EV Certificate**, in Windows 10
- BIOSes are deprecated in favor of UEFI with **Secure Boot**
 - Preventing custom OS/boot loaders from working, even if you have physical access to machine
- Certain processes are **protected** so that not even an **Administrator** can modify them
- On certain platforms, **no 3rd party** user-mode code is **allowed** to run

....and many other mitigations to **prevent Administrators** from **arbitrarily** executing kernel-mode code without TPM measurement, ELAM authorization, SHA-256 EV Code Signing from Microsoft Root CA

And then you **find a way** to get arbitrary kernel-mode code execution which bypasses every single protection put into place

- **Is it a bug?** If not, then why all these mechanisms? Only admins can load drivers. All these mechanisms are meant to **prevent admins** from doing that.

The “Cross-Session Globals”

Defined by a structure which contains 3 main storage containers

Network Font Table and Semaphore Lock

- Used whenever a font image file is being loaded from a remote (UNC) share
- Triggered when the kernel font mapper parses fonts or the font cache
- Couldn't find obvious way to trigger from user mode

Trusted Font Table and Semaphore Lock

- Used whenever a **font** is being **added** to the current process' resource table
- **Inserted** into when the undocumented “**Trusted Font**” flag is used, **queried** when the flag is not used
- Easy to trigger both paths by using **NtGdiAddFontResourceW** exported as **GdiAddFontResourceExW**

Session Pool Tracker Table

- Off by default, can be enabled with undocumented registry key
- Not really exploitable, but can leak addresses if turned on

Windows AVL Tables

Windows implements a “Run-time Library”, or Rtl which provides, among other things, an **AVL Tree/Table Package**

- Self-balancing binary search tree using the AVL algorithm

An **AVL Table** is described by the following structure:

```
lkd> dt nt!_RTL_AVL_TABLE
+0x000 BalancedRoot      : _RTL_BALANCED_LINKS
+0x020 OrderedPointer    : Ptr64 Void
+0x028 WhichOrderedElement : Uint4B
+0x02c NumberGenericTableElements : Uint4B
+0x030 DepthOfTree       : Uint4B
+0x038 RestartKey        : Ptr64 _RTL_BALANCED_LINKS
+0x040 DeleteCount       : Uint4B
+0x048 CompareRoutine    : Ptr64      _RTL_GENERIC_COMPARE_RESULTS
+0x050 AllocateRoutine   : Ptr64      void*
+0x058 FreeRoutine       : Ptr64      void
+0x060 TableContext      : Ptr64 Void
```


Overwriting Compare Callback

With RW access to the section object, we can **replace** the Trusted Font Table with our **own** table

- Ours will contain a **custom compare** callback function, which we control
- The kernel's AVL Table functions will **call** it at **Ring 0** as soon as we try to **load** a font
- WIN!

Except for **SMEP (Supervisor Mode Execution Prevention)**

- Expressly designed to handle lazy weaponization of similar exploits
- CPU will **not allow execution** of CPL 0 code stored in CPL 3 pages (#PF raised)

We need to find a way around SMEP on modern processors

- And traditional way of using **ROP Gadgets** won't help here, as we don't have **stack control**
- **Heap-spraying** or heap **control** works nicely on Windows 7 or x86 Windows 8 by leveraging Executable Big Pool (NonPaged/SessionPaged), but **not x64**.

Convert CodeExec into W-w-w

The “lazy way” is often the “wrong way” with SMEP.

We need to convert our code execution into a **Write-what-where**, and we don’t have ROP gadgets.

Instead, we use an **existing** function to perform a controlled write

- Can make things more complicated and use multiple functions that are called in order by re-exploiting the bug
- In some rare cases, can even use gadgets through **relative offsets (“JOP”)**
 - However, **CFG** will **protect** against that on **Windows 10**

Difficulty is that AVL Compare Routine has a specific set of **parameters**, and must return a **specific** set of possible **return** values

- Function must not only perform the required write, but also return correctly

KASLR is not an issue here: kernel base is **leaked**, and kernel file can be mapped in user-mode

Finding the “right” Compare

The **compare routine** has the following signature

```
RTL_GENERIC_COMPARE_RESULTS  
(*PRTL_AVL_COMPARE_ROUTINE) (  
    __in struct _RTL_AVL_TABLE  *Table,  
    __in PVOID  FirstStruct,  
    __in PVOID  SecondStruct  
);
```

We control **Table**, and we control **SecondStruct**

- How? Compare routine always starts at the **tree root** (controlled in **Table**) and always first starts by taking its right child (i.e.: **Table->Root.RightChild**)
- Returning **2+** means “**Node Found**”, and the **FirstStruct** will be returned.
- Returning **1** means “walk tree down on the **right side**”, so **LeftChild** will be dereferenced as the next **SecondStruct**, and function is called again
- Returning **0** means “walk tree down on the **left side**”, so **opposite** of above

FirstStruct is, in this case, the string containing the font name

Making things hard...

So Table must have **NumberOfGenericElements** set to 1, and **BalancedRoot.RightChild** must be set to an interesting value.

Preferably, then, the function does something like:

- **mov [SecondStruct], [Table->SomeField]**

Which really expands into

- **mov [Table->Root.RightChild.UserData], [Table->SomeField]**
- i.e.: it's important to make sure that **SomeField** is not already being used to maintain the tree consistent.

Return value is important too: if the function returns **2+**, we are golden, because **win32k!IsTrustedFontPath** only checks for **!=NULL**, doesn't actually read the data

- But if the function returns **1**, we must make **Table->Root.LeftChild == NULL**
- If the function returns **0**, we have a problem, because it will read **RightChild**

Best we can get...

```
.text:00000000140160F24 ; char __cdecl PopPepStartDevicePowerOnActivity
.text:00000000140160F24 PopPepStartDevicePowerOnActivity proc near
.text:00000000140160F24 ; DATA
.text:00000000140160F24 mov dword ptr [r8], 2
.text:00000000140160F2B mov rax, [rcx+20h]
.text:00000000140160F2F mov byte ptr [r8+10h], 1
.text:00000000140160F34 mov [r8+8], rax
.text:00000000140160F38 mov al, 1
.text:00000000140160F3A retn
.text:00000000140160F3A PopPepStartDevicePowerOnActivity endp
```

Reads **Table->OrderedPointer** into RAX

Writes RAX into **SecondStruct+8** (or really **Table->Root.RightChild.UserData + 8**)

Unfortunately, corrupts **SecondStruct+0** and **SecondStruct+16**

Dealing with Corruption

MWR (first?) publicized a technique leveraging the fact that in Windows, it is **trivial** to find the **paging hierarchy structures** for a given virtual address

- Fixed **PXE_BASE**, **PPE_BASE**, **PDE_BASE**, **PTE_BASE** constants are well known
- The offset be deduced by looking at the virtual address bits
- In fact this is how **!pte** works!

On x64, Windows 8.1 offers **128 TB** of address space, and few processes ever use anywhere close to this

- Meaning that most of the PXE/PML4 entries are **empty** (each PML4 entry covers **512GB**)
- Therefore, by picking an address “**far enough**” in the address space, and making the target of the **Write-what-where** a given **PML4 entry**, nearby **corruption** won’t matter as long as nobody attempts to dereference anything within the **adjacent 512GB** blocks
- **Repeat** the process for other entries (PPE, PDE, PTE)

Improving the Technique

The original technique uses the **Write-what-where** **four** times for each of the paging structures.

- It describes how **self-referencing** entries work, but **doesn't** actually **leverage** their power!

If we pick a **specialty crafted** virtual address, we can make it such that the **PTE** Offset in the Page Table will be **equal** to the **PDE** Offset in the Page Directory, which itself will be **equal** to the **PPE** Offset in the Page Directory Pointer Table, which itself will be **equal** to the **PXE** Offset in the PML4.

For example, **0x40201008000** uses the **universal** offset **0x40**

```
lkd> !pte 0x40201008000
                                VA 0000040201008000
PXE at FFFFF6FB7DBED040    PPE at FFFFF6FB7DA08040    PDE at FFFFF6FB41008040    PTE at FFFFF68201008040
```

Now only a **single write** is needed, and the PXE will **recursively** be a PTE

Forging the PTE

In the MWR blog post, they allocate an existing PTE, and flip the required **User** and **NoExecute** bits on/off in order to make the page **Supervisor** and **Executable**

- By using a **misaligned** write, they can flip the NX bit specifically off, while writing the required ON bits at the end of the PTE, with a **single write** that **doesn't corrupt** the page frame number (PFN).

Unfortunately, the particulars of our memory corruption make a **misaligned** address **impossible**

- The writes at **+0x08** and **+0x10** will become **misaligned** writes into our PTE, and **corrupt** it

We can **only** write an 8-byte **aligned** value

- So we have to **forge** a **PTE** from **scratch**, including the **correct PFN** and **Working Set Index**

How do we do that?

Address Windowing Extensions

Designed for Windows **2000**, which supported **PAE (Physical Address Extension)**

- Intel technology allowed for access of up to **64GB** of RAM on a **32-bit OS**

Virtual address space remained **2GB**, so how could processes map/access the additional RAM?

- API was created to allow **manual** control of **free RAM** pages
- Used by **SQL Server**

Still used **today**, even on 64-bit due to its useful side-effects

- Pages are not in the **working set**, meaning they don't get **paged/cached** or subject to **trimming** (and faster memory operations are allowed)
- Process can fully determine page management (allows for a user-mode memory manager)

AllocateUserPhysicalPages, MapUserPhysicalPages & MEM_PHYSICAL

How AWE Helps

AllocateUserPhysicalPages returns a “**UserPfnArray**”

UserPfnArray [out]

A pointer to an array to store the page frame numbers of the allocated memory.

The size of the array that is allocated should be at least the *NumberOfPages* times the size of the **ULONG_PTR** data type.

Do not attempt to modify this buffer. It contains operating system data, and corruption could be catastrophic. The information in the buffer is not useful to an application.

“MSDN, you just invited Alex Ionescu to take your API apart”

Bottom line: by **leaking** the address of the **PFNs**, and by **zeroing** out the **working set index**, **AWE** pages’ **PTEs** can be fully **forged** with the needed bits

Protected Process Bypass DEMO

There's One More Thing...

It seems everyone **misunderstands** how **SMEP actually** works

And **AWE** has **one more bug**, which made the entire last 20 minutes useless*

* Except on Windows 10, where it's **partially fixed**

What MWR & Others Got Wrong

*“When checking the rights of a page, the kernel will check every PXE involved in the address translation. That means that if we want to check if the U/S flag is set, we will check all entries relating to that page. **If any of the entries do not have the supervisor flag set**, any attempt to use this page from kernel mode will trigger a fault. **If all of the entries have the supervisor flag set**, the page will be considered a kernel-mode page.”*

What Intel told the NSA...



Supervisor Mode Execution Protection

Stephen Fischer
Senior Principal Engineer
Intel® Corporation

Intel Doublespeak



SMEP architectural control details

- CR4.SMEP – If 1 and in supervisor mode ($CPL < 3$), instructions may not be executed from a linear address for which the U/S flag is 1 (user mode) in every paging-structure entry controlling the translation for the linear address
 - SMEP U/S paging attribute precedence:
 - Any page level marked as supervisor (U/S=0) will result in treatment as supervisor for SMEP enforcement
 - Existing user/supervisor privileging checking continues to require the more conservative mapping (i.e. execution in user mode ($CPL=3$) requires all levels to be mapped as U/S=1 (user))

Just one P*E structure is needed

Because of this, it means that even in a **512GB** address range which is user-mode, and within a **1GB** block that is user-mode, and a **2MB** region that is user-mode, a **single “Supervisor” 4KB** page will **bypass SMEP**

This makes **Write-what-where** vulnerabilities combined with PTE editing an even better deal

- If we didn't have memory corruption (i.e.: a more standard **Write-what-where**), a single edit at a known **PTE_BASE+PteIndex** offset would be enough to **bypass** SMEP.
- Can easily be done **remotely** as **PTE_BASE** is **known**

But it turns out, thanks to **AWE**, you **don't** even **need** a **vulnerability**.

AWE pages can have two states:

- **PAGE_READWRITE**: Non-executable, user-mode, RW
- **PAGE_NOACCESS**: How can we **emulate** “no access”?!

PAGE_NOACCESS on AWE Pages

Normally, **emulating PAGE_NOACCESS** means turning the **Valid** bit off in the **PTE**, and using some **Software PTE** bits, combined with the **VAD**, for the memory manager to understand this is a “real” allocation, but marked as **inaccessible**.

- But **AWE** pages are “**invisible**” or “**unmanaged**” by the memory manager, and it **isn’t easy** to replicate this logic without major changes

So a simpler trick is used instead...

- The kernel maps the page **supervisor**
- This enforces “**no access**” as far as **user-mode** is concerned
- But it also means one of your user page is now **ring 0**!

Even better, the internal array that determines which Win32 permissions (**PAGE_READWRITE** vs **PAGE_EXECUTE_READWRITE**) should flip the **NX** bit **on**, do not include **PAGE_NOACCESS**.

- Kernel will flip the **NX** bit **off**. You get **supervisor, executable** memory.

QUESTIONS?

THANK YOU!